Software Requirements Specification

for

YouTube Video Intelligence MCP Server

Version 1.0

Prepared by Random Stuff

Aadyant Neog BT2024186
Jeswin Jose BT2024142
Paul T Mathew BT2024135

9 November 2026

Copyright © 1999 by Karl E. Wiegers. Permission is granted to use, modify, and distribute this document.

Table of Contents

Table of Contents ii
Revision History iii
1. Introduction 1
1.1 Purpose 1
1.2 Document Conventions 1
1.3 Intended Audience and Reading Suggestions 1
1.4 Product Scope 1
1.5 References 2
2. Overall Description 3
2.1 Product Perspective 3
2.2 Product Functions 3
2.3 User Classes and Characteristics 3
2.4 Operating Environment 4
2.5 Design and Implementation Constraints 4
2.6 User Documentation 4
2.7 Assumptions and Dependencies 4
3. External Interface Requirements 6
3.1 User Interfaces 6
3.2 Hardware Interfaces 6
3.3 Software Interfaces 6
3.4 Communications Interfaces 7
4. System Features 7
4.1 Similar Videos 7
4.2 Similar Channels 8
4.3 Personalised Recommendations 8
4.4 Engagement Analysis 9
4.5 Constrained Playlist Creation 10
4.6 Authorisation and Account Management 12
4.7 Quota Budgeting 12
4.8 Input Validation Coverage and Errors 13
5. Other Nonfunctional Requirements 14
5.1 Performance Requirements 14
5.2 Safety Requirements 14
5.3 Security Requirements 15
5.4 Software Quality Attributes 15
5.5 Business Rules 16
6. Other Requirements 16
6.1 Agent Skill Requirements 16
6.2 Engineering and Knowledge Requirements 17
6.3 Verification and Traceability 18
Appendix A: Glossary 21
Appendix B: Analysis Models 22
Appendix C: To Be Determined List 24

Revision History

Name	Date	Reason For Changes	Version
Random Stuff	7 November 2026	Initial YouTube SRS and team roster.	1.0
Random Stuff	9 November 2026	Reorganised into the teaching assistants' template; requirement statements retained.	1.0 IEEE

1. Introduction

1.1 Purpose

This Software Requirements Specification defines version 1.0 of the YouTube Video Intelligence MCP Server for Team Random Stuff. It covers the complete project baseline: the five required product features, supporting interfaces, quality and security constraints, agent skills, engineering obligations, and acceptance evidence. The topic and submitted requirements remain fixed after submission. This document specifies required behaviour and does not certify completed implementation or human approval.

1.2 Document Conventions

Shall denotes a mandatory, testable obligation. All numbered requirements are mandatory unless a stated condition applies. FR denotes functional requirements, NFR nonfunctional requirements, SK agent-skill requirements, GOV engineering obligations, and AT planned acceptance scenarios. The original requirement IDs and statements are retained. All system features have High priority because the supplied assignment requires them; priorities do not make detailed requirements optional. References [R1] to [R12] retain their original meanings.

1.3 Intended Audience and Reading Suggestions

The instructor and reviewers should read Sections 1 and 2 for scope, Section 4 for behaviour, Section 5 for quality, and Appendix C for unresolved dependencies. Developers should additionally study Section 3, the skill and engineering requirements in Section 6, and Appendix B. Testers should use each requirement's verification link and the acceptance catalogue in Section 6. Team members are responsible for understanding the complete submitted baseline.

1.4 Product Scope

Learners use YouTube to study specific topics but must manually sift through distracting results, compare creators, check runtimes, and assemble playlists. The proposed system shall let a learner use a conversational AI client to discover relevant videos and channels, understand their accessible YouTube library, and create a focused playlist within explicit time and language constraints.

The product is a Model Context Protocol server connected to the official YouTube Data API v3. It shall provide metadata-based similarity, personalised recommendations, engagement summaries, and a preview-and-confirm playlist workflow. The server shall validate constraints and enforce approval before writing to the user account. Reusable agent skills and a linked Second Brain shall support development and demonstration.

The baseline excludes scraping, unofficial APIs, video or audio downloading, arbitrary third-party transcript retrieval, watch-history access, a replacement YouTube player, uploads, comments, likes or subscription changes, and editing or deleting existing playlists. A custom web application, public hosted service, model training, and support for every AI client are not required. The client may summarise API metadata but shall not claim to have watched a video or read unavailable captions.

1.5 References

The supplied project statement and common SDLC document govern scope. The following technical references were consulted for the baseline on 7 November 2026; this template adaptation does not assert a new external-source review. The structural template was supplied by the teaching assistants as srs_template-ieee.doc, based on the Karl E. Wiegers template, copyright 1999, whose permission notice is retained.

[R1] Project Statement YouTube Video Intelligence MCP Server. Supplied project brief, Project_1_YouTube_MCP.pdf, pages 1-2. Authority for FR-1 to FR-5, video-to-playlist, TQ-1 to TQ-6 and official-API constraints.

[R2] The AI Assisted Software Development Life Cycle. Supplied common workflow, SDLC_and_Workflow.pdf, pages 1-3. Authority for Second Brain, engineering, skills, CI, human review and stage deliverables.

[R3] YouTube Data API reference. Google for Developers. Official resource and method catalogue. https://developers.google.com/youtube/data/reference/rest

[R4] Videos list. Google for Developers. Video metadata, authorised myRating retrieval and result limits. https://developers.google.com/youtube/data/reference/rest/v3/videos/list

[R5] Playlist items and subscriptions. Google for Developers. Playlist-addition and subscription-creation timestamp semantics. See also the subscriptions resource in the same reference. https://developers.google.com/youtube/data/reference/rest/v3/playlistItems

[R6] Quota calculator. Google for Developers. Current bucket allocations, endpoint costs and reset boundary; profile checked 7 November 2026. https://developers.google.com/youtube/data/determine_quota_cost

[R7] OAuth 2 for desktop apps. Google for Developers. Installed-app consent, loopback callbacks and PKCE. https://developers.google.com/identity/protocols/oauth2/native-app

[R8] Using OAuth 2 to access Google APIs. Google for Developers. Refresh-token expiry, revocation and external Testing-mode seven-day lifetime. https://developers.google.com/identity/protocols/oauth2

[R9] Playlist creation and item insertion. Google for Developers. Authorised write methods and scopes; see also playlistItems.insert in the same reference. https://developers.google.com/youtube/data/reference/rest/v3/playlists/insert

[R10] Model Context Protocol tools specification. Official MCP specification, revision 2025-11-25. Tool schemas, discovery, invocation and human interaction guidance; accompanying resources and transports sections define their interfaces. https://modelcontextprotocol.io/specification/2025-11-25/server/tools

[R11] Agent Skills specification. Official skill directory, SKILL.md frontmatter and instruction format. https://agentskills.io/specification

[R12] YouTube API Services Developer Policies. Google for Developers. Storage and deletion requirements and derived-data restrictions; see Section III E 4 and dependency D-02. https://developers.google.com/youtube/terms/developer-policies

2. Overall Description

2.1 Product Perspective

This is a new local integration component used by an existing conversational AI client. It is not a replacement for YouTube or the AI model. The server owns validation, quota tracking, metadata processing, proposal state and account-write safeguards; external services own video hosting, account identity and consent. Appendix B shows the system context and playlist lifecycle.

The learner interacts with an existing AI client. The client calls the local MCP server, which validates requests, queries a YouTube provider, computes results, manages proposals and approvals, and returns structured responses. The provider alone calls the official YouTube Data API over HTTPS. Google handles user sign-in and consent. The AI model and YouTube are external dependencies.

2.2 Product Functions

Feature	**Required outcome**
FR-1 Similar videos	Rank videos related to a supplied video URL or ID.
FR-2 Similar channels	Rank channels related to a channel URL, handle, or ID.
FR-3 Recommendations	Use authorised likes, subscriptions, and playlists as preference signals.
FR-4 Engagement analysis	Report top channels and categories, duration preferences, and supported time comparisons.
FR-5 Playlist creation	Interpret a request, propose a compliant playlist, obtain approval, and create it.

2.3 User Classes and Characteristics

User class	**Characteristics and permitted work**
Learner	Primary user. Supplies a seed/topic and constraints, grants Google consent, reads results and approves or rejects a proposed playlist. Technical API knowledge is not required.
Team developer	Configures local credentials and dependencies, maintains provider integrations, skills, tests and brain notes. Must understand code and documentation submitted under the team workflow.

User class	Characteristics and permitted work
Instructor or reviewer	Examines requirements and evidence, observes demonstrations and checks traceability. Does not receive the learner's credentials or unrestricted account access.

2.4 Operating Environment

The baseline shall run locally for one authorised account per server instance on a documented Windows development environment. The implementation shall include a reproducible install or build command and dependency versions. A supported desktop MCP client and a browser for Google consent are prerequisites. Internet connectivity, an enabled Google Cloud project, credentials, and a YouTube channel are required for live account operations. Fixture mode shall work without live credentials.

The performance reference is a machine with at least four CPU cores and 8 GB RAM, as defined in NFR-08. No GPU, embedded device or special peripheral is required. The supported client/runtime release and dependency versions are recorded during design and installation without changing the five feature commitments.

2.5 Design and Implementation Constraints

Use only the official YouTube Data API v3, with access isolated behind a replaceable provider interface. Compute similarity from accessible metadata; do not assume watch history or related-video/channel endpoints. Keep tests independent of live API access and do not create extra Cloud projects to evade quota. Account mutations are restricted to creating approved new playlists and their items. The local stdio boundary, OAuth constraints, security controls and quota profile are specified in Sections 3 to 5.

2.6 User Documentation

Deliver a setup guide, client-connection guide, consent/privacy explanation, feature examples, playlist approval and recovery instructions, and known limits. Developer documentation shall describe MCP tools/resources, schemas, provider interfaces and test execution. Keep these as versioned Markdown and client-readable help resources. If a REST interface is later introduced, its API documentation is governed by GOV-04; no REST interface is required by this baseline.

2.7 Assumptions and Dependencies

ID	Condition and required response
A-01 Client boundary	At least one supported client can load skill instructions, display structured results and provide trusted confirmation. A documented local approval path shall be used if the client cannot provide trustworthy approval.

ID	Condition and required response
A-02 API data limits	Only accessible, budget-bounded metadata is available. Unknown language/duration, missing tags, API result caps and absent historical events reduce coverage and shall be disclosed.
A-03 Consent and channel	The user supplies consent, configured test-user access and a YouTube channel for creation. Denial, Testing-mode expiry or missing channel shall produce guided recovery.
D-01 Quota changes	The handout describes the 10,000-unit default. Current official documentation additionally separates search into its own daily allowance. The implementation shall use a dated configurable quota profile rather than a permanently hard-coded assumption.
D-02 API policy review	The assignment requires computed similarity and aggregated analysis. YouTube Developer Policies Section III E 4 restricts derived data and metrics. The team shall resolve the applicability of those restrictions with the instructor and, where needed, the API provider before enabling affected live processing or public deployment. Labelling an application score is not itself evidence of permission.
R-01 Data and content quality	Metadata can be incomplete, misleading or changed after preview. Show its evidence and retrieval time; revalidate before creation. Do not claim verified educational quality.
R-02 Partial remote writes	YouTube operations are separate API calls, so playlist creation and insertion are not one atomic transaction. Preserve operation state, disclose partial results and reconcile uncertain outcomes.
R-03 Evaluation evidence	Use fixtures to verify required algorithms and error paths while external dependencies are unavailable. Keep fixture and live evidence distinct. An unresolved live dependency remains a block to the affected live acceptance scenario.

Appendix C tracks the unresolved policy and design-selection items. They do not silently remove core features or establish that live acceptance has passed.

3. External Interface Requirements

3.1 User Interfaces

The existing AI client provides the conversational interface. Users enter video/channel identifiers or natural-language goals and receive readable results with source links, durations and explanations. The proposal view shows the account, privacy, exact ordered videos, constraints and warnings. A trusted client control or local human interaction supplies approval; a tool argument alone is insufficient. Detailed visual layout depends on the selected client and is documented during design.

3.2 Hardware Interfaces

The product uses standard operating-system process, disk, credential-storage and network facilities on the development machine. It has no direct interface to sensors, actuators or specialist hardware. Browser-based consent and keyboard/mouse interaction are supplied by the host environment.

3.3 Software Interfaces

NFR-01 MCP interoperability. The server shall support the standard local stdio transport, negotiated protocol initialisation, tool discovery and invocation, and resource discovery and reading. Tool inputs shall have machine-readable schemas; outputs shall include machine-readable fields and a readable explanation. Diagnostic logging shall use stderr and shall not corrupt stdout protocol messages. At least one documented MCP client shall be demonstrated.

Verification: AT-01 and AT-19; [R10].

Interface	**Contract**
Conversational client	Accept natural-language requests; display ranked results, warnings, proposals and approval controls.
Google OAuth	Browser-based installed-app authorisation with a loopback callback, state validation and PKCE. Tokens do not pass through chat.
YouTube Data API v3	Retrieve video, channel, subscription and playlist metadata; create only an approved new playlist and its items.
MCP resources	Expose non-secret capability/help text, current quota status and account-scoped proposal/report resources.
Agent skills	Ship versioned SKILL.md workflows. Make instructions discoverable as resources and loadable by the selected skill-capable client.

The YouTube provider exchanges validated identifiers and structured resource data. Account tokens remain in protected local storage. Proposal, quota and operation records are account-scoped. Appendix B defines the required data objects. Agent Skills remain versioned SKILL.md instruction packages; exposing instructions as resources does not turn skills into a new MCP protocol primitive.

3.4 Communications Interfaces

The client launches the MCP server as a local process and exchanges structured protocol messages over stdio. Diagnostics use stderr. Calls to Google OAuth and the YouTube Data API use HTTPS. The OAuth callback binds to loopback only, and the product does not require a public listener, email, FTP or a custom browser service. Schema validation, safe error responses, per-attempt timeouts and account isolation apply at these boundaries. See NFR-02, NFR-07, NFR-09 and FR-8.

4. System Features

Each feature below is required. Descriptions and response sequences explain the observable workflow; the numbered requirements are the binding details. Error cases and shared controls apply across features. Acceptance tests are planned scenarios, not claimed results.

4.1 Similar Videos

4.1.1 Description and Priority

Priority: High. Given a seed video, identify distinct videos with explainable metadata overlap.

4.1.2 Stimulus/Response Sequences

1. User or client: Provide a supported video URL or ID and a result limit. System: Validate and normalise it, preflight the request and retrieve available seed/candidate metadata.

2. User or client: Request the ranked results. System: Compute similarity, exclude the seed, remove duplicates and return stable ordering with evidence; report an unavailable seed or insufficient candidates explicitly.

4.1.3 Functional Requirements

FR-1.1 Retrieve a seed and candidate pool. Given a valid video URL or ID, the server shall retrieve its accessible metadata and discover a bounded pool of candidates through official API searches or channel metadata. An unavailable seed shall produce a typed error. Candidate retrieval shall not assume a related-videos endpoint exists.

Verification: AT-02 and AT-10; [R1, R3, R4].

FR-1.2 Compute and explain similarity. The server shall compute video similarity locally from available titles, descriptions, tags and categories using an explainable technique such as TF-IDF with cosine similarity or documented weighted scoring. It shall return at most the requested number of unique candidates in descending score order, exclude the seed, and break ties by a documented stable rule. Each result shall include its URL, title, channel, duration, application score and short metadata-based reason.

Verification: AT-02; unit tests for scores, ordering and ties.

FR-1.3 Handle incomplete evidence. Missing tags or descriptions shall not crash ranking. The server shall use remaining supported features and disclose reduced evidence. It shall return an empty result with a reason when no eligible candidates are available. Results shall not describe similarity as verified educational quality or as a recommendation issued by YouTube.

Verification: AT-02 and AT-10.

4.2 Similar Channels

4.2.1 Description and Priority

Priority: High. Resolve a channel and compare its metadata and bounded video sample with other channels.

4.2.2 Stimulus/Response Sequences

1. User or client: Provide a channel URL, handle or ID. System: Resolve the seed and discover comparable channels within budget.

2. User or client: Inspect the comparison. System: Return distinct channel rankings, selection reasons and sample limitations, or a corrective error.

4.2.3 Functional Requirements

FR-2.1 Resolve and compare channels. Given a channel URL, @handle or ID, the server shall resolve the seed channel and rank distinct candidate channels by available channel descriptions and a bounded sample of their video metadata. The response shall disclose the sample size and selection rule. It shall exclude the seed channel and shall not rely on an assumed related-channels endpoint.

Verification: AT-03; [R1, R3].

FR-2.2 Return explainable channel rankings. Each channel result shall include ID, title, URL, application score and a short explanation of the topic overlap. Identical inputs, fixtures and configuration shall produce identical ordering. Unresolvable handles, inaccessible channels or insufficient comparable metadata shall return an explicit error or insufficient-data result.

Verification: AT-03 and AT-10.

4.3 Personalised Recommendations

4.3.1 Description and Priority

Priority: High. Use accessible, authorised likes, subscriptions and playlists as preference evidence.

4.3.2 Stimulus/Response Sequences

1. User or client: Authorise read access and request recommendations, optionally with filters. System: Retrieve the available collections, report coverage and build the profile.

2. User or client: Inspect recommended content. System: Show evidence-linked rankings. If all usable signals are absent, request a seed or topic instead of inventing preferences.

4.3.3 Functional Requirements

FR-3.1 Read authorised preference signals. With read consent, the server shall retrieve accessible liked videos, subscriptions, and playlists and their accessible items. It shall form a preference profile from the available sources, deduplicate video IDs across collections, and identify each source and its coverage. It shall not obtain or infer an actual watch-history feed.

Verification: AT-04; [R1, R4].

FR-3.2 Recommend and explain. The server shall discover and rank a bounded set of videos using that profile and any explicit topic, duration or language filters. Each result shall identify the signals supporting the recommendation. Results shall exclude seed/profile videos by default and shall disclose this exclusion. The same fixtures and settings shall produce a stable ranking.

Verification: AT-04 and AT-15.

FR-3.3 Support sparse and denied data. If a source is empty or inaccessible, the server shall report the limitation and use remaining consented sources. If no usable preference data exists, it shall request a seed video or topic rather than invent preferences. Permission errors shall never be bypassed by scraping or another account.

Verification: AT-04 and AT-11.

4.4 Engagement Analysis

4.4.1 Description and Priority

Priority: High. Summarise accessible collection activity and compare supported time periods without claiming watch history.

4.4.2 Stimulus/Response Sequences

1. User or client: Select collections and optional non-overlapping periods. System: Retrieve supported records and distinguish playlist-addition, subscription-creation and unavailable event times.

2. User or client: Request the report. System: Return channel/category counts, duration statistics and valid period comparisons with denominators, missing data and date bases.

4.4.3 Functional Requirements

FR-4.1 Summarise accessible collections. The server shall report the top channels and video categories by item count, duration preferences through count, mean and median video length, and counts in documented duration bands. It shall distinguish liked videos, subscriptions and playlist items, identify the denominator of each count or percentage, and state duplicate handling, missing metadata and the retrieved population. Subscription summaries shall count channels rather than pretend subscriptions are video views.

Verification: AT-05 with hand-calculated fixtures.

FR-4.2 Compare periods with valid timestamps. The server shall support two user-specified non-overlapping periods for trends and comparisons using available playlist-addition and subscription-creation timestamps. UTC intervals shall be start-inclusive and end-exclusive. Reports shall name the timestamp basis and include absolute differences; percentage change from a zero baseline shall be marked undefined. Video publication dates shall not be substituted for engagement dates. Unsupported like-event dates shall be reported as unavailable.

Verification: AT-05, including boundaries, zero baselines and missing dates; [R5].

FR-4.3 Avoid unsupported conclusions. Reports shall label the data as accessible collection activity and shall not claim watch time, completion, full historical behaviour or learning progress. A video runtime describes the video, not time actually watched. Empty periods shall return zero supported counts and a clear insufficient-evidence explanation for other measures.

Verification: AT-05; [R1].

4.5 Constrained Playlist Creation

4.5.1 Description and Priority

Priority: High. Interpret a playlist request, show a feasible proposal, obtain exact human approval and create the playlist safely.

4.5.2 Stimulus/Response Sequences

1. User or client: Specify topic, total duration, maximum item length and language. System: Clarify ambiguity and produce a bounded, compliant proposal or an infeasibility explanation.

2. User or client: Inspect, revise, reject or approve the proposal. System: Make no writes on rejection; invalidate earlier approval after a change. Accept approval only through the trusted human path.

3. User or client: Approve the current account-bound version. System: Recheck scope, channel, constraints and quota, then create exactly the approved playlist and items.

4. User or client: Inspect success or failure. System: Return the actual URL and completion state. Reconcile uncertain writes and disclose partial results without duplicate creation.

4.5.3 Functional Requirements

FR-5.1 Translate and validate a request. The AI client and server together shall accept a natural-language playlist request specifying topic, total duration, maximum video length and language. The client shall translate these into structured values; the server shall independently validate them. Missing or ambiguous hard constraints shall be clarified before proposing a playlist. A request for an exact total shall be clarified as an upper limit or rejected as unsupported, rather than silently reinterpreted.

Verification: AT-06 and AT-10.

FR-5.2 Enforce hard constraints. The server shall build an ordered proposal whose total duration does not exceed the requested total and whose individual durations do not exceed the requested maximum. Every selected video shall satisfy the stated topic-selection rule and requested language according to available API language metadata. Unknown language shall not count as a strict match, and search relevanceLanguage alone shall not prove spoken language. Unverifiable candidates shall be excluded and the evidence limitation shall be shown.

Verification: AT-06, AT-07 and duration-boundary unit tests.

FR-5.3 Handle insufficient candidates. The server shall report if no non-empty feasible proposal can be built. If only a shorter compliant playlist is available, it shall show the total and unused time and allow the user to accept it or revise the request. It shall never relax a hard constraint, add duplicate filler or claim exact duration without evidence.

Verification: AT-07.

FR-5.4 Present the complete proposal. Before any playlist write, the client shall display the destination account/channel, title, privacy, topic and constraints, ordered video titles and links, channels, per-video duration, total duration, language evidence, warnings and proposal expiry. Default privacy shall be private. The proposal shall have an account-bound ID and version tied to the exact content. Any change to account, videos, order, title, privacy or constraints shall invalidate earlier approval.

Verification: AT-06, AT-08 and AT-12.

FR-5.5 Require trusted human approval. The write operation shall require explicit human approval of the current proposal through a trusted client confirmation control or a local approval interaction outside model-controlled tool arguments. Approval shall be bound to the user/account, proposal version and operation, expire within 10 minutes, and be usable for only that operation. A supplied confirmed=true value or instructions embedded in metadata shall not establish approval. Rejected, missing, expired, mismatched or superseded approval shall cause zero playlist writes.

Verification: AT-08 and AT-12.

FR-5.6 Create and report the playlist. After valid approval, the server shall verify write scope, destination channel, current feasibility and quota, then create the new playlist and insert exactly the approved videos in the approved order. It shall return the playlist URL, final item count and duration, and a success or partial-failure status. It shall not silently replace an unavailable video or change privacy. A material change found during revalidation shall require a new preview and approval.

Verification: AT-06, AT-09 and AT-11; [R9].

FR-5.7 Prevent duplicate writes and disclose partial state. The server shall record the operation and known remote IDs. Repeating a completed operation ID shall return its recorded result without another playlist creation. On an insertion failure it shall stop and report the playlist URL, inserted and pending items. An ambiguous network timeout on a write shall be reconciled before any retry; if unresolved, it shall report an unknown outcome and require user action. Resume shall target only verified missing items, with renewed approval if needed. Automatic deletion is outside the baseline.

Verification: AT-09.

Worked constraint example: a request for up to two hours with no item longer than 20 minutes means a total of at most 7,200 seconds and an item maximum of 1,200 seconds. A 1,201-second item fails even if rounded display suggests 20 minutes. A feasible 6,420-second proposal shows 107 minutes used and 13 unused. Unknown language/duration does not count as proof of compliance.

4.6 Authorisation and Account Management

4.6.1 Description and Priority

Priority: High. Establish delegated account access and recover from consent expiry, missing scope or missing YouTube channel.

4.6.2 Stimulus/Response Sequences

1. User or client: Start a workflow requiring private data or a write. System: Obtain the minimum sufficient Google consent through the browser and keep tokens outside chat.

2. User or client: Disconnect or encounter revoked/expired access. System: Clear local account data on disconnect, attempt revocation and provide specific recovery guidance when access fails.

4.6.3 Functional Requirements

FR-6 Manage consent and account state. The server shall support authorisation, access-token refresh, re-authorisation after expiry or revocation, and disconnect. Read workflows shall request youtube.readonly; creation shall request a documented sufficient write scope such as youtube.force-ssl only when needed. The client shall explain scope purpose. Missing channel or youtubeSignupRequired shall return guidance without claiming creation succeeded. Disconnect shall remove local tokens and account-specific stored data and attempt token revocation.

Verification: AT-11 and AT-14; [R7, R8, R9].

4.7 Quota Budgeting

4.7.1 Description and Priority

Priority: High. Reject unaffordable workflows before provider traffic and account for bounded calls safely.

4.7.2 Stimulus/Response Sequences

1. User or client: Request any API-backed workflow. System: Estimate and reserve the cost of the complete bounded operation across applicable buckets.

2. User or client: Continue execution or exhaust the available balance. System: Track attempts and release unused reservations; reject insufficient budgets before any provider call.

4.7.3 Functional Requirements

FR-7 Preflight and account for quota. Before the first provider call in a workflow, the server shall estimate and reserve its bounded worst-case cost across all applicable quota buckets, including pagination, revalidation, writes and permitted retries. If any bucket is insufficient, it shall reject the workflow before any provider call. It shall track each attempted call, including errors, and release unused reservations. Parallel operations shall not oversubscribe the tracked balance.

Verification: AT-13 and AT-15; [R6].

4.8 Input Validation Coverage and Errors

4.8.1 Description and Priority

Priority: High. Apply shared input and coverage rules and return safe, actionable error states.

4.8.2 Stimulus/Response Sequences

1. User or client: Supply an identifier, limit, date range or constraint. System: Validate the schema and normalise supported forms before calls.

2. User or client: Receive limited or failed results. System: Identify pagination, truncation and inaccessible content; return stable error codes with a safe next action and no secrets.

4.8.3 Functional Requirements

FR-0.1 Normalise and validate inputs. The server shall accept canonical YouTube video URLs, youtu.be links and video IDs, and channel URLs, @handles and channel IDs. It shall validate identifiers, positive durations, result limits, language codes, date ranges, title and privacy fields before provider calls. Unknown hosts or ambiguous legacy channel URLs shall return corrective guidance instead of being fetched as arbitrary URLs.

Verification: AT-02, AT-03 and AT-10.

FR-0.2 Declare coverage and limits. Every collection result shall state whether it is complete, paginated, sampled or truncated and shall report inaccessible items. The baseline shall support result limits from 1 to 50, a configurable candidate cap with a default of 200 unique videos, and proposals containing 1 to 50 videos. Larger requests shall return a clear limit message. Library pagination shall stop at the configured request budget and report the continuation state; it shall not label partial data as a full account history.

Verification: AT-04 and AT-15.

FR-8 Return recoverable errors. The server shall return stable error codes, a readable cause, retryability and a safe next action for invalid input, unavailable content, insufficient data, consent failure, quota exhaustion, rate limiting and provider failures. Errors shall not leak credentials, private metadata or stack traces. Partial results shall be explicitly marked.

Verification: AT-10, AT-11 and AT-16.

5. Other Nonfunctional Requirements

5.1 Performance Requirements

NFR-08 Meet local performance targets. On a documented reference machine with at least four CPU cores and 8 GB RAM, fixed offline fixtures and no client-model generation time, input validation shall have a 95th-percentile latency of at most 100 ms; ranking up to 200 candidate videos or constructing a proposal shall take at most 2 seconds. Tests shall use at least 30 measured runs after one warm-up. Cached read responses shall have a 95th percentile of at most 2 seconds. Live network and OAuth latency shall be reported separately.

Verification: AT-17 benchmark report with hardware, inputs and timings.

NFR-09 Bound runtime and resource use. A provider HTTP attempt shall have a configurable timeout with a default of 15 seconds. A foreground read or planning workflow shall stop after a configurable deadline with a default of 60 seconds, return a typed timeout or labelled partial result, and perform no writes. Creation shall report progress and honour the same per-attempt timeout; a stop shall preserve partial/unknown state. All searches and pagination shall respect the declared candidate and quota caps.

Verification: AT-15 and AT-16 with simulated delays and bounded-call assertions.

NFR-12 Track changing API quotas. Quota method costs, bucket limits and daily reset rules shall be configurable and versioned with an authoritative reference date. The initial profile shall account for a 10,000-unit daily general bucket and the separately documented 100-call daily search allowance. The local ledger shall survive server restart and use the provider's reset boundary. It shall expose estimated remaining budget and warn that other applications sharing the Cloud project can consume unobserved quota.

Verification: AT-13 and AT-15; [R6].

Performance thresholds apply to the documented local fixture environment. Network availability, OAuth interaction and model generation are external latency factors and must be reported separately. The quota profile remains dated and configurable rather than treating the handout's simplified quota summary as immutable.

5.2 Safety Requirements

The product does not control physical equipment. Its principal loss risks are unintended account changes, exposed personal data and duplicate or partially completed writes. The safeguards are exact human confirmation (FR-5.5), revalidation and honest completion reporting (FR-5.6), operation reconciliation (FR-5.7), and fail-safe recovery (NFR-06). No safety certification is claimed. Metadata-based recommendations must not be represented as a guarantee of factual or educational quality.

5.3 Security Requirements

NFR-02 Protect credentials and transport. All external requests shall use HTTPS. OAuth access and refresh tokens shall be stored through OS-protected credential storage or an equivalently protected encrypted local store, with the key outside the repository. Secrets shall never appear in source, fixtures, brain notes, chat, logs, resources or tool results. Automated secret scanning shall run on pull requests.

Verification: AT-14; inspect repository, logs and token storage.

NFR-03 Enforce authority boundaries. The server shall validate schemas and account bindings on every tool call and resource read. It shall treat titles, descriptions, tags and retrieved text as data, never instructions. Only the trusted human approval path shall authorise playlist writes; the model and external metadata shall not be able to mint approval. Cross-account or modified-proposal access shall fail.

Verification: AT-12 and AT-14 adversarial tests.

NFR-04 Minimise personal data. Only data needed for the requested workflow shall be collected. Private account data shall remain isolated from other accounts and shall not be sent to the AI client except as needed for the requested result. Fixtures and CI logs shall contain synthetic or fully anonymised data. A user-facing privacy notice shall explain collection, AI-client disclosure, retention, disconnect and local deletion.

Verification: AT-14 and documentation inspection.

NFR-05 Limit retention and support removal. API metadata caches shall expire within 24 hours by default and shall be refreshed or deleted no later than 30 days. Local proposal and operation records containing API data shall follow that limit. Disconnect or a user deletion request shall clear account-specific caches, proposals, operation data and tokens during successful completion. Failed removal shall be reported and shall not be marked complete. Data deletion in this application shall not claim to delete the user's YouTube content.

Verification: AT-14 with simulated time and deletion checks; [R12].

NFR-07 Constrain network destinations. The server shall parse supported YouTube identifiers locally and shall not fetch arbitrary user-supplied URLs. Provider traffic shall be limited to official Google OAuth and YouTube API destinations required by the workflow. The OAuth callback shall bind only to the local loopback interface. The baseline shall not expose a public network listener.

Verification: AT-10 and AT-14; inspect outbound requests and callback binding.

5.4 Software Quality Attributes

Correctness of constraints, explicit approval and protection of user data take precedence over ranking breadth and latency. Reliability, explainability, maintainability and repeatable testing are required as specified below. Interoperability is defined by NFR-01 in Section 3.3.

NFR-06 Fail safely and recover. Read requests may use at most three attempts with bounded exponential backoff and jitter for retryable failures. Quota exhaustion and permanent permission errors shall not trigger repeated calls. Write timeouts shall follow FR-5.7. After interruption, stored operation state shall distinguish not-started, in-progress, completed, partial and unknown outcomes. A failed write shall never be reported as complete.

Verification: AT-09 and AT-16 fault-injection tests.

NFR-10 Provide usable and explainable results. The client shall display human-readable durations, clickable source URLs, collection coverage, relevant constraints and actionable errors. Scores shall be labelled application estimates and shall identify the metadata used. Confirmation shall explicitly identify account and privacy. User and developer documentation shall include setup, consent, each feature, limits, recovery and one complete example.

Verification: AT-19 task walkthrough and documentation review.

NFR-11 Support maintainable testing. YouTube access shall be isolated behind a provider interface replaceable by a fixture provider. Domain rules for similarity, duration parsing, constraints and approval shall be testable independently of a live API or AI model. Dependencies and the supported runtime shall be declared. CI shall build or validate the package, lint, test, scan secrets and validate skills and brain notes on every PR.

Verification: AT-18 and AT-20; [R2].

5.5 Business Rules

Rule	**Enforcing requirements**
Only the authorised learner can approve a specific current playlist proposal; OAuth alone is not that approval.	FR-5.4 to FR-5.6; FR-6
A total-duration request is an upper bound. Do not fill unused time by violating language, topic, duration or uniqueness constraints.	FR-5.1 to FR-5.3
Writes create a new playlist and its approved items only; default privacy is private.	FR-5.4; FR-5.6; scope in Section 1.4
Untrusted metadata cannot grant authority or alter budget rules.	FR-5.5; FR-7; NFR-03
Retain the submitted topic and requirement baseline; changes need instructor direction.	GOV-05

6. Other Requirements

6.1 Agent Skill Requirements

Each skill is a versioned directory with valid SKILL.md name/description frontmatter, invocation conditions, inputs, instructions, outputs and failure handling. Every skill cites a distilled brain note. Client setup must document how instructions are loaded; skills cannot override server validation or approval.

SK-01 brain ingest. The brain-ingest skill shall capture supplied material or research as source-linked notes, record authorship and date, add requirement IDs and related wikilinks, update the index, and append the run to the log. Stage 1 shall contain at least 10 distinct research notes created through this workflow. Unverified material shall be identified as such.

Verification: AT-18: ingest sample material; inspect at least 10 source-linked research notes.

SK-02 brain refine. The brain-refine skill shall organise notes into PARA, merge duplicates while preserving sources and authorship, and repair or report broken links. It shall report the changes made and leave an auditable history through the normal PR workflow.

Verification: AT-18: duplicate-note and broken-link fixtures.

SK-03 brain query. The brain-query skill shall answer project questions or draft a document from relevant notes, cite the notes supporting its claims, and disclose missing or conflicting evidence. It shall not invent decisions, approvals, test results or sources.

Verification: AT-18 and AT-19: answer a requirement-to-decision question and an unsupported question.

SK-04 local check. The local-check skill shall build or validate the project, run the automated test suite, start the local server, invoke one read-only fetch tool through MCP, and report pass, fail or blocked separately for every step. It shall terminate resources it started. Fixture mode shall be the default; an explicitly configured live check shall report quota used.

Verification: AT-19: live demonstration of all five steps with individual statuses.

SK-05 video to playlist. The video-to-playlist skill shall accept a seed video, summarise its available title/description metadata with a stated evidence basis, find similar videos through FR-1, collect playlist constraints, propose a playlist, obtain human confirmation, create it through FR-5, and present a short analysis of duration, channels and constraints. Rejection or an infeasible request shall end without writes. It shall not claim access to full video content or unavailable transcripts.

Verification: AT-06, AT-08 and AT-19; full accepted and rejected workflow evaluations.

SK-06 Evaluate skills before submissions. Each applicable skill shall have repeatable evaluation cases covering a successful path, a negative path and evidence/authority handling. Evaluations shall run before each relevant stage submission and results shall be recorded with inputs, expected behaviour, observed behaviour and pass/fail/blocked status.

Verification: AT-18; stage-specific evaluation records, not generated claims of success.

6.2 Engineering and Knowledge Requirements

GOV-01 Maintain linked project knowledge. The repository shall contain an Obsidian-compatible brain vault with inbox, projects, areas, resources and archive directories, an index entry point and an append-only run log. Notes shall link to other notes, requirement IDs and real issues/PRs as available. Every team member shall author notes under their own identity.

Verification: AT-20; repository and commit-history review; [R2] SB-1, SB-6.

GOV-02 Record decisions and AI work. The repository shall contain root AGENTS.md and significant design decisions in docs/adr/. Each ADR and skill shall cite at least one distilled note. Significant AI sessions shall record goal, outcome and what worked or failed. The team shall write a review note at each stage and shall be able to explain every merged artefact.

Verification: AT-20; [R2] SB-2 to SB-4 and AI policy.

GOV-03 Use reviewed version control. All work and defects shall be tracked as GitHub Issues. Code, skills, brain and document changes shall use meaningful incremental commits and PRs with at least one human review. The team shall document its branching strategy and maintain release tags and a changelog. Contributions shall remain attributable to their actual authors.

Verification: AT-20; inspect issues, commit history, PR reviews and release records; [R2].

GOV-04 Maintain engineering artefacts. The team shall maintain architecture and design documentation with suitable UML/ER diagrams, ADRs, a development plan and task breakdown, user documentation and developer documentation for MCP tools/resources and provider interfaces. An OpenAPI document is required only if a REST interface is introduced.

Verification: AT-20; [R2] SE-2 to SE-4 and SE-7.

GOV-05 Preserve traceability and the submitted baseline. The team shall maintain requirement-to-issue-to-PR-to-test-to-skill-to-brain-note traceability. Before submission, requirement edits shall be reviewed. After submission, the selected topic and SRS shall remain the fixed baseline; implementation decisions shall refine it without silently changing scope. Instructor-directed clarifications shall be recorded separately.

Verification: AT-20; traceability completeness and baseline comparison.

Stage	**Required evidence**
1 Requirements	SRS with FR/NFR/SK; PARA vault, index/log, at least 10 ingested research notes; AGENTS.md; repository and CI skeleton.
2 Design	Architecture, diagrams and ADRs; brain-refine and brain-query; initial local-check.
3 Testing and CI	Continuous testing/CI plan and results; evaluated video-to-playlist; defect log and regression tests; stage reviews.
4 Review	Live FR-1 to FR-5 demonstration; product skill, local-check and brain-query; traced requirement walkthrough; final release tag.

6.3 Verification and Traceability

The following catalogue defines planned acceptance evidence. Each requirement above names the applicable scenarios. Product implementation, test execution, human review and submission are not established by this document.

ID	**Scenario and required evidence**
AT-01	MCP initialisation, tool schemas, resources, stdio cleanliness and selected-client connection.
AT-02	Video ranking: seed exclusion, duplicates, stable ties, missing tags and empty candidates.
AT-03	Channel URL/handle/ID resolution and explainable, distinct channel rankings.

ID	Scenario and required evidence
AT-04	Personalisation using all three signal types, sparse sources and visible pagination.
AT-05	Hand-calculated analysis, UTC boundaries, empty periods and unavailable event times.
AT-06	Successful constrained proposal, approval, exact creation and video-to-playlist flow.
AT-07	Duration/language/topic constraints, unknown metadata, underfill and infeasibility.
AT-08	Reject, absent, expired, replayed, changed or wrong-account approval; zero new writes.
AT-09	Partial creation, restart, ambiguous timeout, reconciliation and duplicate prevention.
AT-10	Malformed IDs, unsupported hosts, invalid limits and private/deleted content.
AT-11	Expired/revoked tokens, denied scope, failed refresh and youtubeSignupRequired.
AT-12	Metadata prompt injection and model-forged approval cannot bypass confirmation.
AT-13	Insufficient general or search bucket: reject before any provider call; concurrent reservation.
AT-14	Secret storage/redaction, least privilege, isolation, local deletion and cache expiry.
AT-15	Pagination/candidate caps, cost accounting, reset/restart and externally exhausted quota.
AT-16	Rate limiting, timeouts and provider 5xx; bounded safe retries and typed failures.
AT-17	Repeatable offline performance measurements against NFR-08.
AT-18	Skill evaluations; notes, links and frontmatter validation; offline CI execution.
AT-19	User-facing walkthroughs, documentation and live skill/client demonstration.
AT-20	Design, ADR, plan, review, traceability and release-evidence inspection.

Source requirement	Required verification
TQ-1 Unit	Test domain functions, services and tool handlers, especially similarity, ISO 8601 duration parsing and constraint boundaries.
TQ-2 Integration	Test MCP handlers, services, state storage and interchangeable YouTube providers together.
TQ-3 System	Run AT-02 to AT-06 through the MCP interface for all five core features.

Source requirement	Required verification
TQ-4 Negative	Cover invalid IDs, private/deleted videos, missing tags, quotaExceeded, youtubeSignupRequired, expired tokens and rate limits.
TQ-5 CI and live separation	PR CI uses synthetic fixtures/mocks with no live credentials. A separate scheduled live workflow has a fixed, configurable quota budget and authorised test account.
TQ-6 Quota	Test full-workflow budget preflight before any API call, accounting for both current bucket profiles and configured cost changes.
TQ-7 Skills	Evaluate applicable skills before every stage submission and retain the results.
TQ-8 Regression	Maintain tests as the product changes and give every significant defect a regression case.

Requirements	Acceptance	Skill	Brain note
FR-0.1 to FR-0.2	AT-02/03/10/15	local-check	04, 10
FR-1.1 to FR-1.3	AT-02/10	video-to-playlist	02
FR-2.1 to FR-2.2	AT-03/10	local-check	02
FR-3.1 to FR-3.3	AT-04/11/15	local-check	03
FR-4.1 to FR-4.3	AT-05	local-check	05
FR-5.1 to FR-5.7	AT-06 to AT-12	video-to-playlist	04, 06, 07
FR-6 to FR-8	AT-09 to AT-16	local-check	07, 08, 09, 10
NFR-01 to NFR-12	AT-01/09 to AT-20	local-check	07 to 12
SK-01 to SK-06	AT-06/08/18/19	Each named skill	11, 12
GOV-01 to GOV-05	AT-20	brain skills	12

Tests are planned, not executed results. Acceptance shall require every applicable mandatory requirement to have passing evidence, with unresolved failures and external blocks explicitly listed. Security-critical failures involving approval bypass, secret exposure, cross-account access or duplicate writes shall block release. Missing live credentials or a policy decision shall be reported as blocked, never passed. Regression cases shall be added for significant defects.

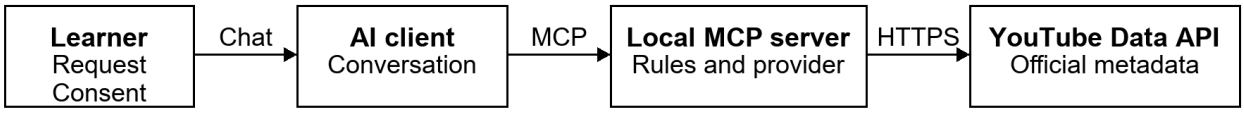
Brain note numbers identify the numbered source notes in the repository index. Actual issue, PR and implementation-test links are added as work is created; missing links are pending rather than fabricated. The companion IEEE requirement-placement map records the new section for each unchanged baseline requirement.

Appendix A: Glossary

Term	Definition
MCP	Model Context Protocol; the interface by which a client discovers and calls server tools and reads resources.
SRS	Software Requirements Specification; the project's required behaviour and constraints.
OAuth	Delegated authorisation to an account; separate from confirmation of a particular playlist.
TF-IDF	Term-frequency inverse-document-frequency; a text feature representation for similarity.
Cosine similarity	Comparison of the directions of two feature vectors.
ADR	Architecture Decision Record, linked to supporting knowledge notes.
CI	Continuous integration checks run on proposed repository changes.
PARA	Projects, Areas, Resources and Archive; the organising system for the Second Brain.
Fixture	Synthetic or fully sanitised test data used without live account access.
Proposal	An account-bound read-only playlist plan that still requires human approval.
Quota bucket	A provider-defined allowance against which particular API methods are accounted.
TBD	A tracked item awaiting a decision; it does not make mandatory features optional.

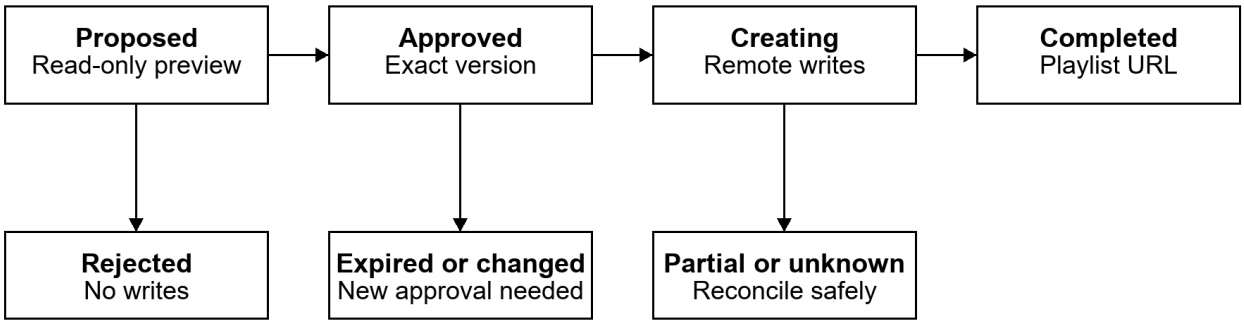
Appendix B: Analysis Models

These conceptual models explain the system boundary and required behaviour. They do not prescribe an implementation language, database product or detailed class design.



The server checks account consent, constraints, quota and human approval.
Results return through the same interfaces. Google handles browser sign-in.

Figure B1 System context



Replayed completed operations return the recorded result without duplicate creation.

Figure B2 Playlist approval and execution lifecycle

In Figure B2, a changed or expired proposal returns to the preview/approval workflow. Only successful revalidation permits writes. Partial or unknown outcomes require reconciliation; completed-operation replays return the recorded result. See FR-5.4 to FR-5.7.

Core information objects and their required fields follow. Unknown metadata remains explicitly unknown, and all account-specific objects require the account binding specified by the baseline.

Object	**Required information**
Video	ID, title, channel ID/title, canonical URL, description, available tags/category, publication time, duration in seconds, language evidence, fetch time and availability. Missing fields remain explicitly unknown.
Channel	ID, handle when available, title, canonical URL, description and the metadata sample used for comparison.

Object	**Required information**
Library collection	Authorised account binding, collection type, accessible items, retrieval time, pagination/completeness and skipped-item counts.
Ranked result	Video or channel identifier, rank, application-computed score, selection reason and evidence limitations.
Playlist proposal	Proposal ID/version, account binding, constraints, title, privacy, ordered unique video IDs, durations, total, evidence, expiry and approval state.
Execution result	Operation ID, proposal version, created playlist ID/URL, inserted and remaining items, status and safe recovery advice.

Appendix C: To Be Determined List

The following items are tracked for closure. They preserve the existing policy dependency and leave implementation selections to design; they do not permit silent changes to the submitted topic or mandatory features.

ID	Item and closure evidence	Owner and timing
TBD-01	Resolve dependency D-02: applicability of YouTube derived-data/metric restrictions to the assigned similarity and analytics features. Record the instructor/provider interpretation and permitted live scope in an ADR. Until then, distinguish fixture evidence from blocked live acceptance.	Project team with instructor/provider; before affected live use.
TBD-02	Select and document the client/version and trusted human confirmation route required by NFR-01 and FR-5.5. Verify connection, resources and exact-proposal approval.	Project team; design and client-integration stage.
TBD-03	Record the implementation runtime, dependency versions and ranking configuration under NFR-11. These selections must meet the preserved requirements and the documented local environment.	Project team; design and reproducible-build setup.